EXPERIMENT NO. 04

Aim:

To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy Your First Kubernetes Application.

Theory:

Kubernetes is an open-source platform for deploying and managing containers. It provides a container runtime, container orchestration, container-centric infrastructure orchestration, self- healing mechanisms, service discovery and load balancing. It's used for the deployment, scaling, management, and composition of application containers across clusters of hosts.

But Kubernetes is more than just a container orchestrator. It could be thought of as the operating system for cloud-native applications in the sense that it's the platform that applications run on, just as desktop applications run on MacOS, Windows, or Linux.

From a high level, a Kubernetes environment consists of a control plane (master), a distributed storage system for keeping the cluster state consistent (etcd), and a number of cluster nodes.

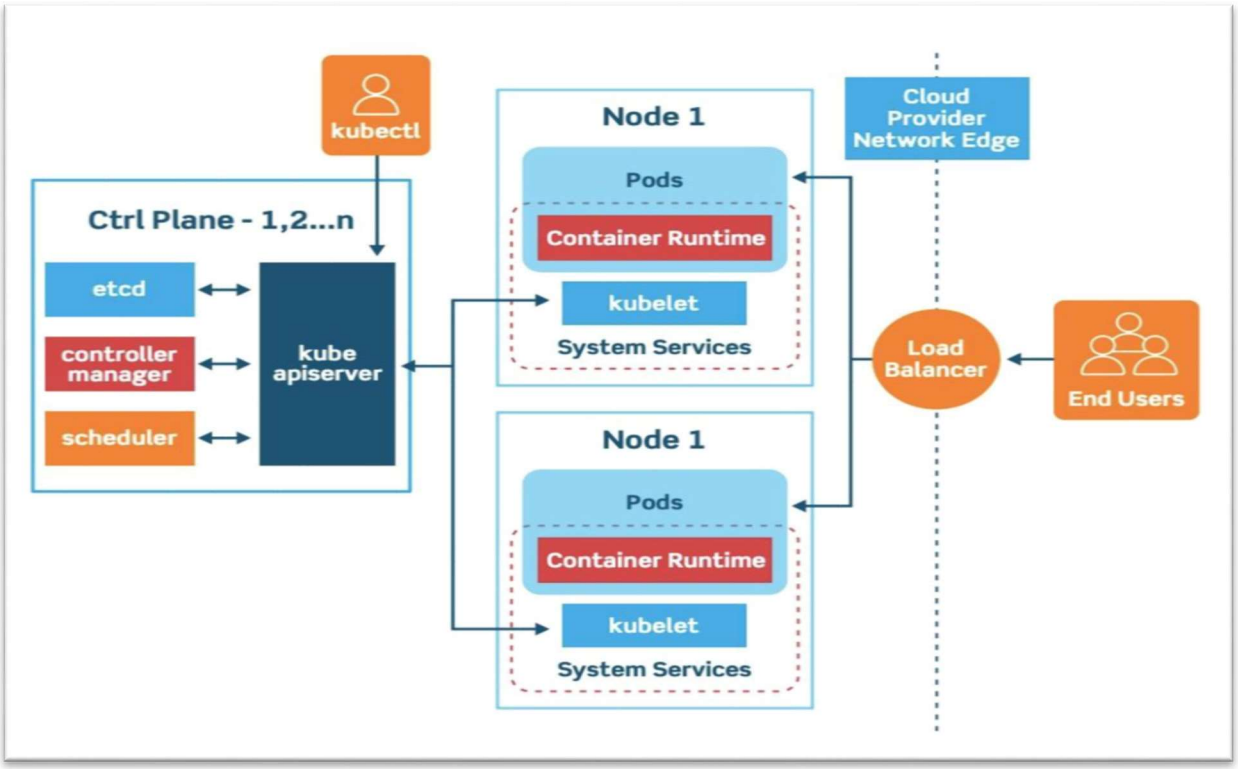

The **control plane** is the system that maintains a record of all Kubernetes objects. It continuously manages object states, responding to changes in the cluster; it also works to make the actual state of system objects match the desired state. As the above illustration shows, the control plane is made up of three major components:

- kube-apiserver
- kube-controller-manager
- kube-scheduler

Cluster nodes are machines that run containers and are managed by the master nodes. The Kubelet is the primary and most important controller in Kubernetes. It's responsible for driving the container execution layer, typically Docker.

Pods are one of the crucial concepts in Kubernetes, as they are the key construct that developers interact with. The previous concepts are infrastructure-focused and internal architecture

Kubernetes Tooling and Clients:

Here are the basic tools you should know:

1. **Kubeadm** bootstraps a cluster. It's designed to be a simple way for new users to build clusters (more detail on this is in a later chapter).
2. **Kubectl** is a tool for interacting with your existing cluster.
3. **Minikube** is a tool that makes it easy to run Kubernetes locally. For Mac users, HomeBrew makes using Minikube even simpler.

- **KUBERNETES DEPLOYMENT**

A Kubernetes deployment is a resource object in Kubernetes that provides declarative updates to applications. A deployment allows you to describe an application's life cycle, such as which images to use for the app, the number of pods there should be, and the way in which they should be updated.

A Kubernetes object is a way to tell the Kubernetes system how you want your cluster's workload to look. After an object has been created, the cluster works to ensure that the object exists, maintaining the desired state of your Kubernetes cluster.

The process of manually updating containerized applications can be time consuming and tedious. Upgrading a service to the next version requires starting the new version of the pod, stopping the old version of a pod, waiting and verifying that the new version has launched successfully, and sometimes rolling it all back to a previous version in the case of failure.

Performing these steps manually can lead to human errors, and scripting properly can require a significant amount of effort, both of which can turn the release process into a bottleneck.

A Kubernetes deployment makes this process automated and repeatable. Deployments are entirely managed by the Kubernetes backend, and the whole update process is performed on the server side without client interaction.

A deployment ensures the desired number of pods are running and available at all times. The update process is also wholly recorded, and versioned with options to pause, continue, and roll back to previous versions.

Automate deployments with pre-made, repeatable Kubernetes patterns

The Kubernetes deployment object lets you:

- Deploy a replica set or pod
- Update pods and replica sets
- Rollback to previous deployment versions
- Scale a deployment
- Pause or continue a deployment

Pre-requisite:

- Launch Linux server running Ubuntu 18.04 /20.04 on Virtual box

OR

- Launch EC2 instances of Ubuntu 20.04 AMI free tier.

Steps:

Step 1: To create nginx run the command kubectl create deployment nginx –image=nginx

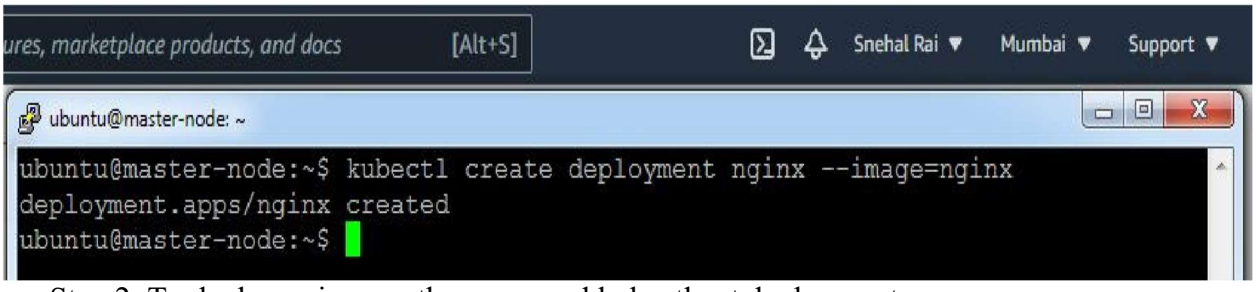

Step 2: To deploy nginx run the command kubectl get deployment

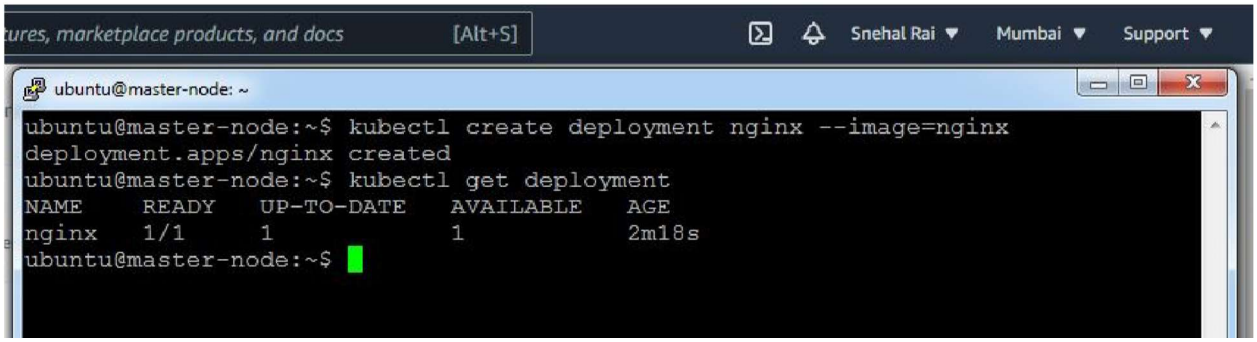

Step 3: Run the command kubectl expose deploy nginx –-port 80 –target-port 80 –type NodePort

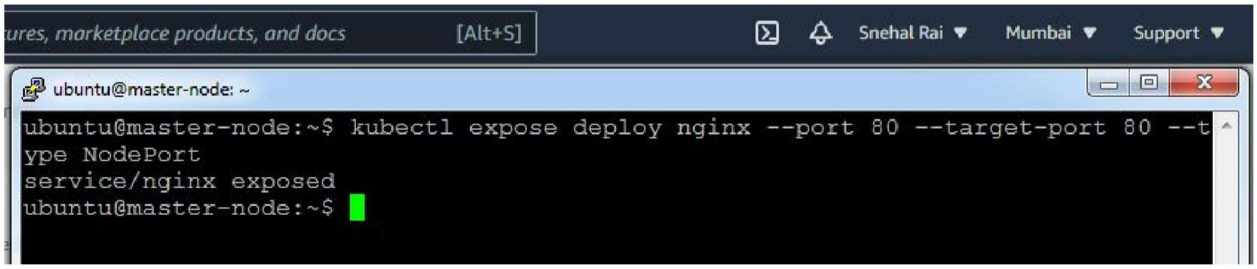

Step 4: To know thw service run the command kubectl get service

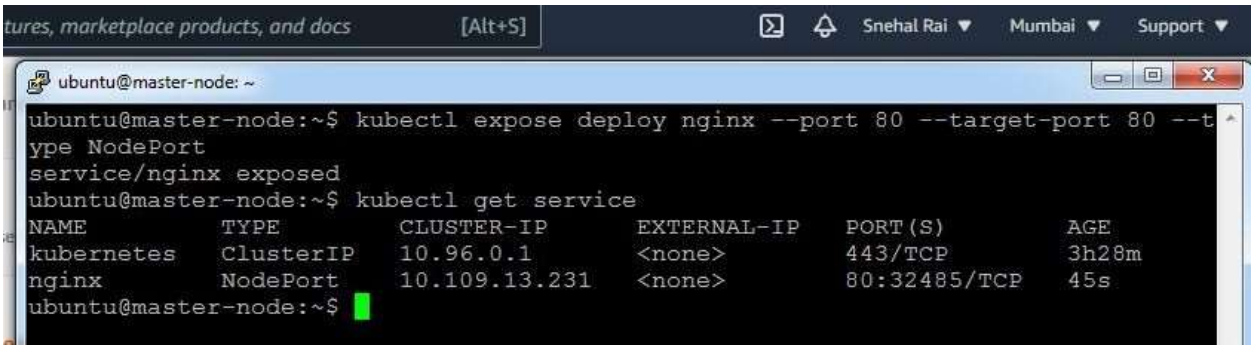

Step 5: Nginx will open in the browser

Step 6: Run the command kubectl get pods

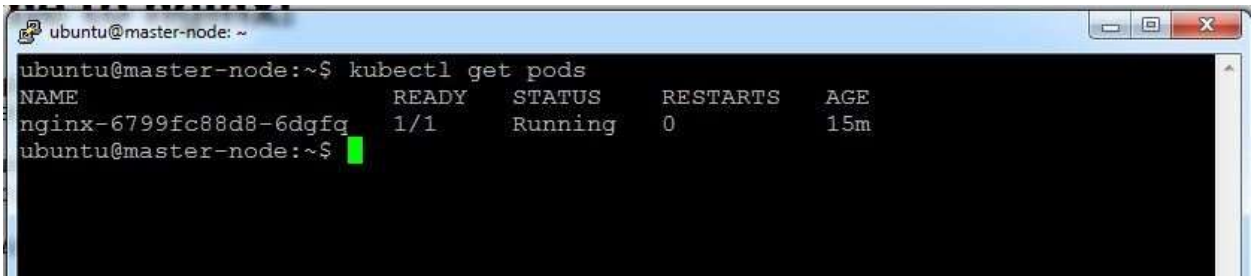

Step 7: To scale nginx run the command kubectl scale ---current-replicas=2 deployment/nginx

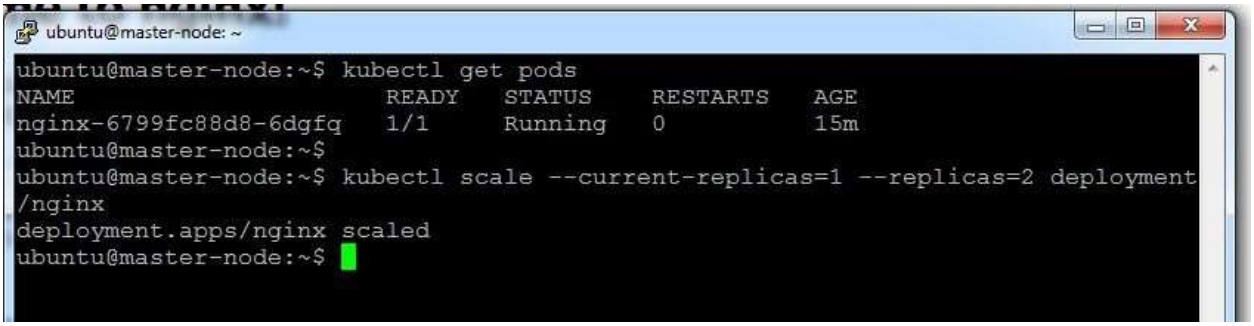

Step 8: To see the status again run the command kubectl get pods

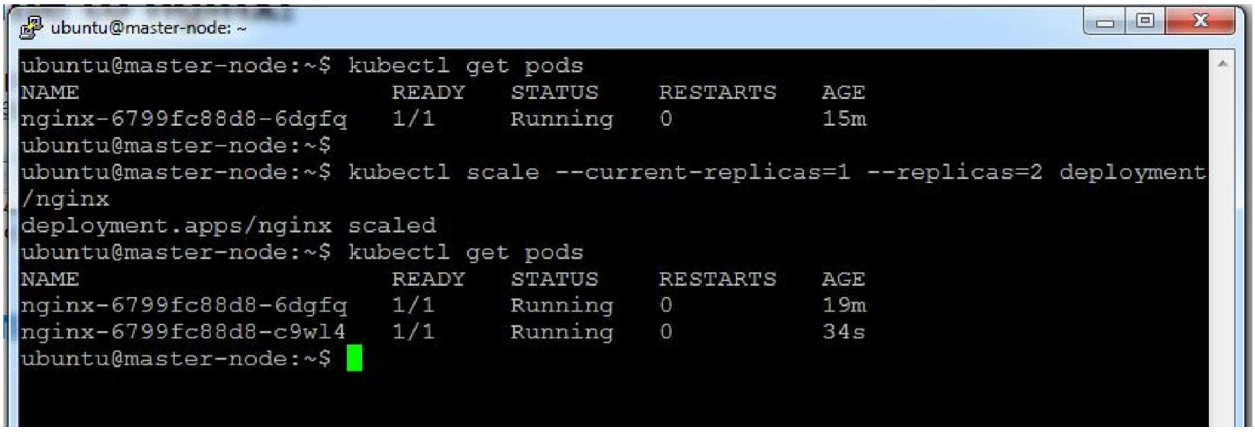

Step 9: For nginx pod description run the command kubectl describe pods nginx –6799fc88d86dgfq

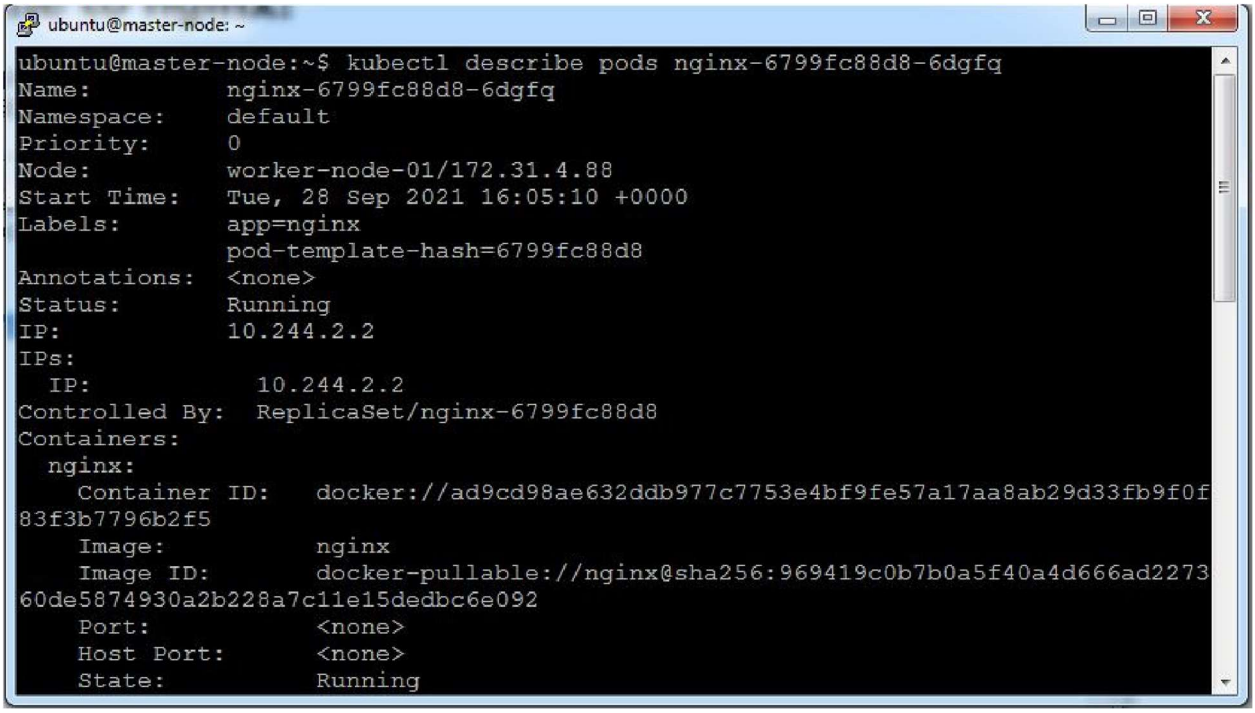

Step 10: For nginx deployment description run the command kubectl describe deployment/nginx

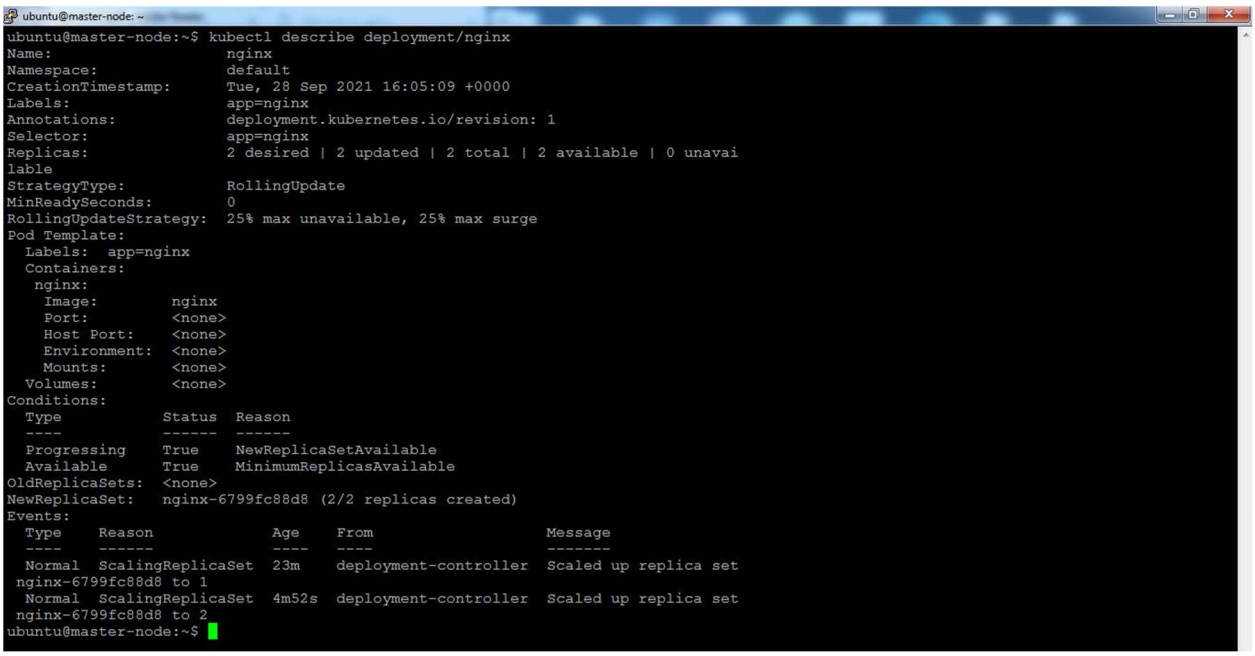

Step 11: Run the command kubectl get services

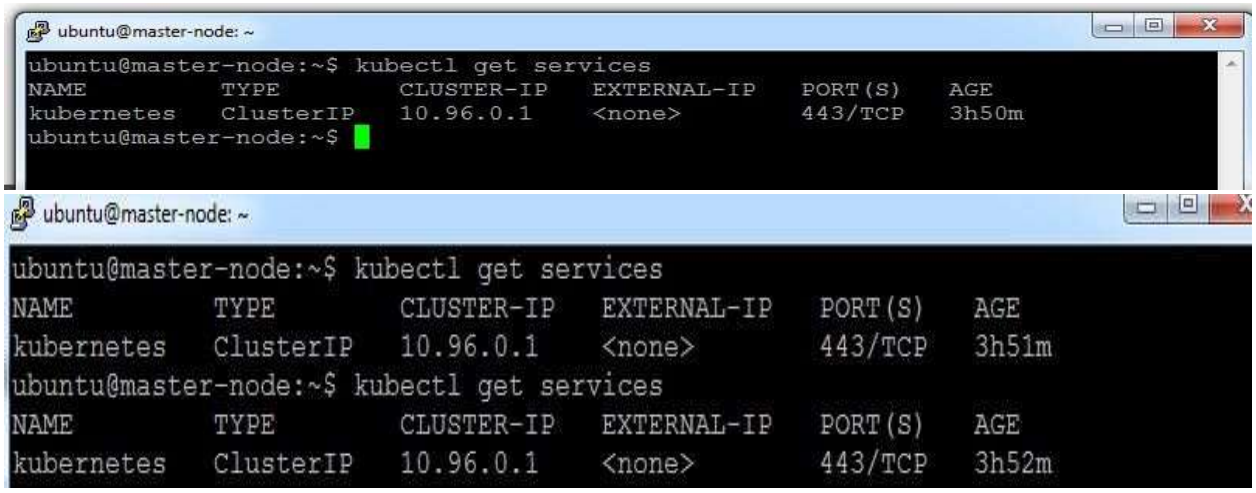

Step 12 : To delete the deployment of nginx run the command kubectl delete deployment nginx

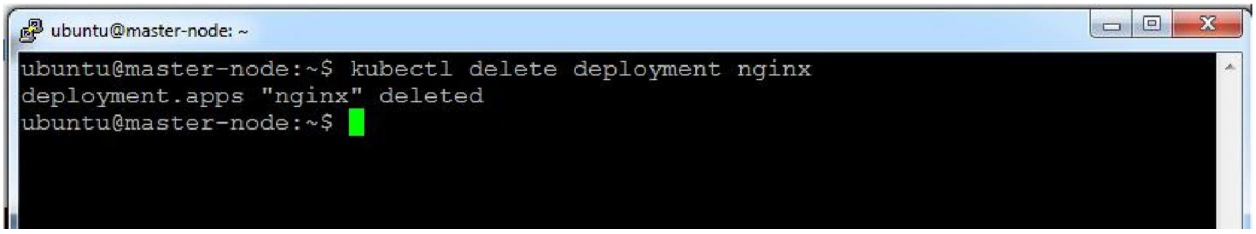

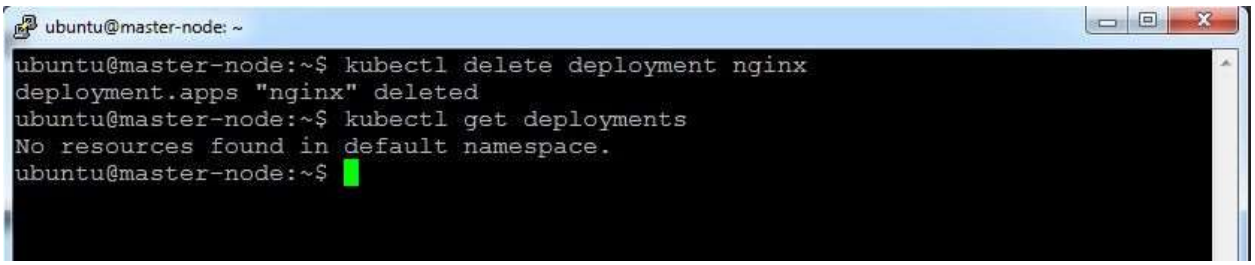

Conclusion:

Kubernetes is an open-source container orchestration platform. It provides a complete platform for scaling and managing applications that are deployed in containers. NGINX provides a suite of products which run within Kubernetes environments: Continuous

development and integration, the rapid deployment and elasticity of containers and cloud services, and breaking our applications into interconnected microservices are emerging as the new normal.NGINX perform HTTP routing, directing each request to the appropriate server as defined by policies that refer to values in the Host header and URI, followed up by load balancing, health checks, and session persistence.