EXPERIMENT 3

Aim:

To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

Theory:

Kubernetes is an open-source platform for deploying and managing containers. It provides a container runtime, container orchestration, container-centric infrastructure orchestration, self- healing mechanisms, service discovery and load balancing. It's used for the deployment, scaling, management, and composition of application containers across clusters of hosts.

But Kubernetes is more than just a container orchestrator. It could be thought of as the operating system for cloud-native applications in the sense that it's the platform that applications run on, just as desktop applications run on MacOS, Windows, or Linux.

From a high level, a Kubernetes environment consists of a control plane (master), a distributed storage system for keeping the cluster state consistent (etcd), and a number of cluster nodes.

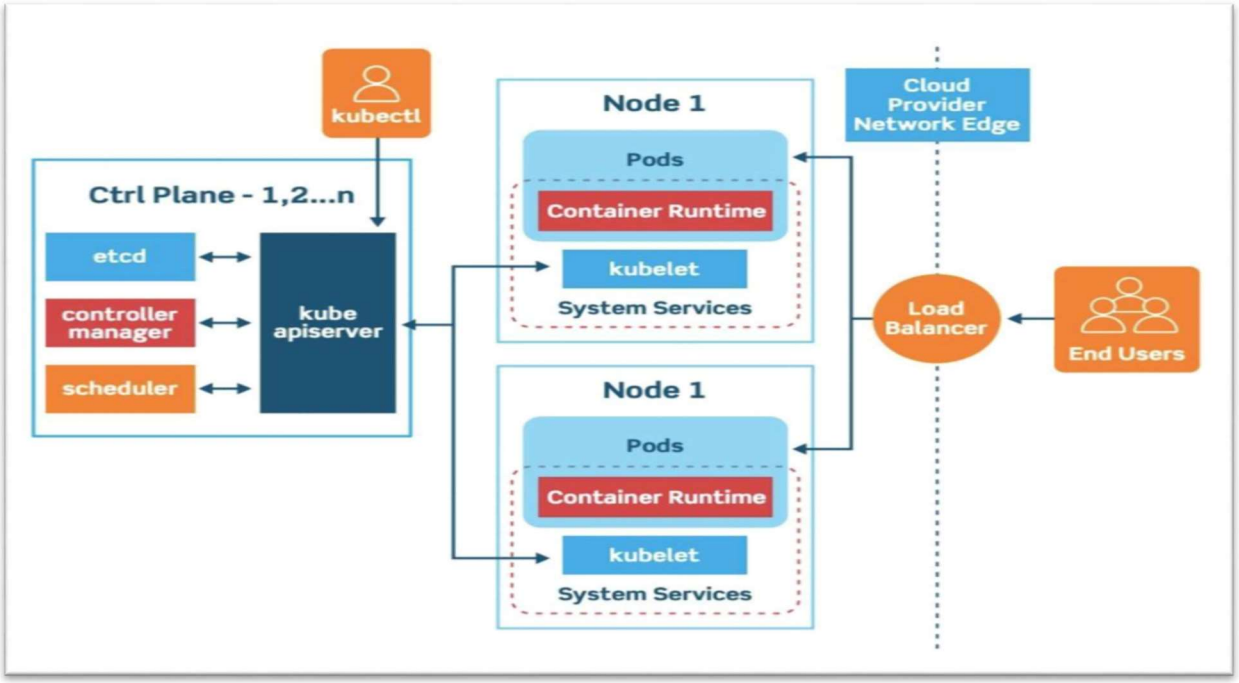

The control plane is the system that maintains a record of all Kubernetes objects. It continuously manages object states, responding to changes in the cluster; it also works to make the actual state of system objects match the desired state. As the above illustration shows, the control plane is made up of three major components:

- kube-apiserver
- kube-controller-manager
- kube-scheduler

Cluster nodes are machines that run containers and are managed by the master nodes. The Kubelet is the primary and most important controller in Kubernetes. It's responsible for driving the container execution layer, typically Docker.

Pods are one of the crucial concepts in Kubernetes, as they are the key construct that developers interact with. The previous concepts are infrastructure-focused and internal architecture

Kubernetes Tooling and Clients:

Here are the basic tools you should know:

1. Kubeadm bootstraps a cluster. It's designed to be a simple way for new users to build clusters (more detail on this is in a later chapter).
2. Kubectl is a tool for interacting with your existing cluster.
3. Minikube is a tool that makes it easy to run Kubernetes locally. For Mac users, HomeBrew makes using Minikube even simpler.

Pre-requisite:

- Launch two or more Linux servers running Ubuntu 18.04 /20.04 on Virtual box

  OR

- Launch two or more EC2 instances of Ubuntu 20.04 AMI free tier.

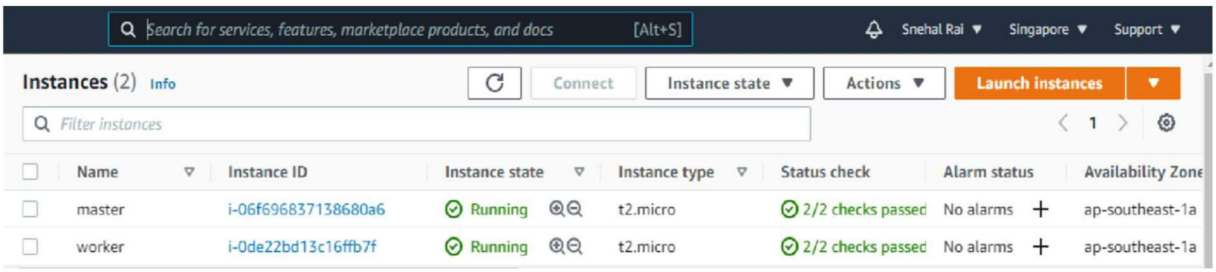

- If using EC2 instances, connect to all instances using PUTTY (on Windows) or using.

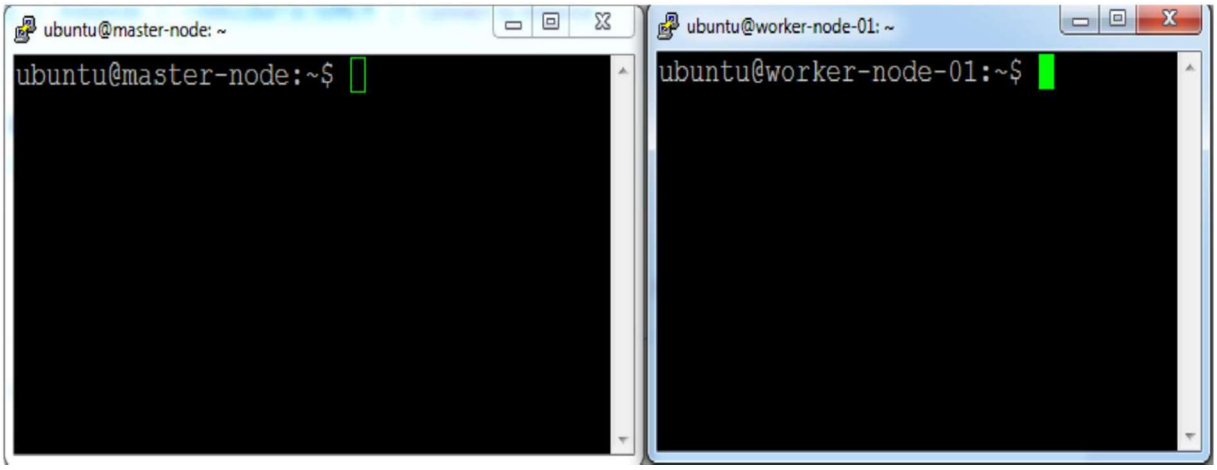

- Access to a user account on each system with sudo or root privileges.
- The apt package manager is included by default.

Steps to Install Kubernetes on Ubuntu.

- Set up Docker

Step 1: Install Docker

Kubernetes requires an existing Docker installation. If you already have Docker installed, skip ahead to Step 2. If you do not have Kubernetes, install it by following these steps:

Update the package list with the command:

$ sudo apt-get update

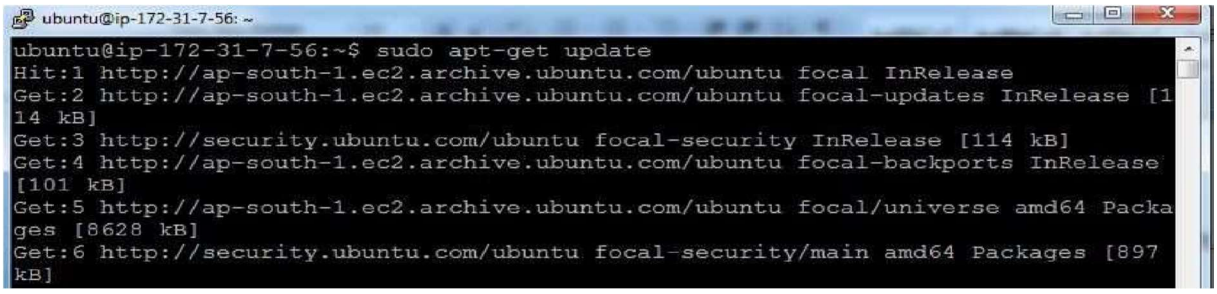

Next, install Docker with the command:

$ sudo apt-get install docker.io

Repeat the process on each server that will act as a node.

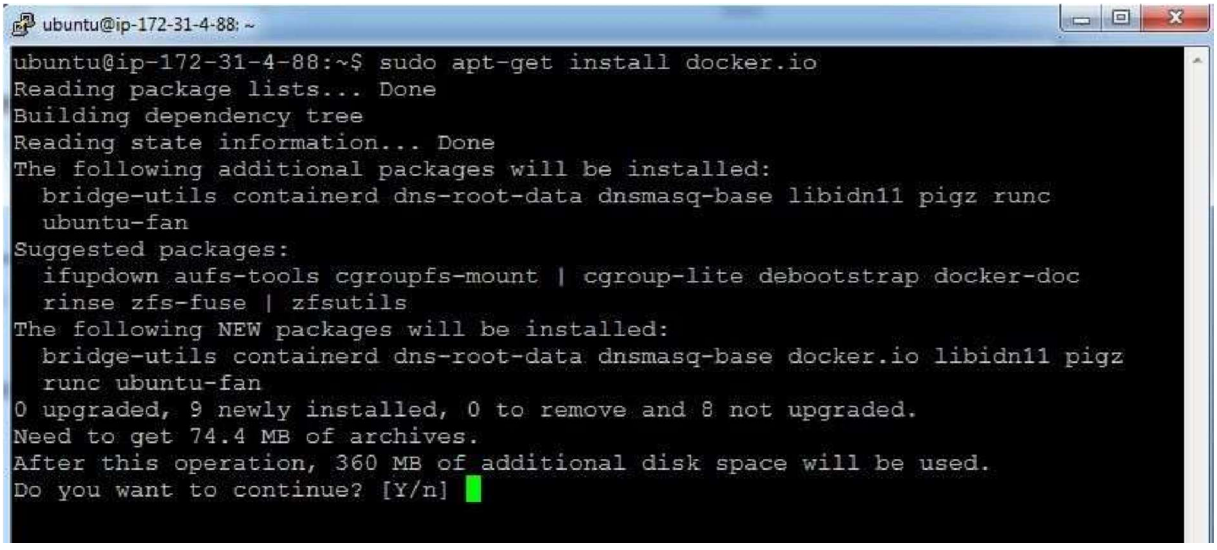

Check the installation (and version) by entering the following:

$ docker —version

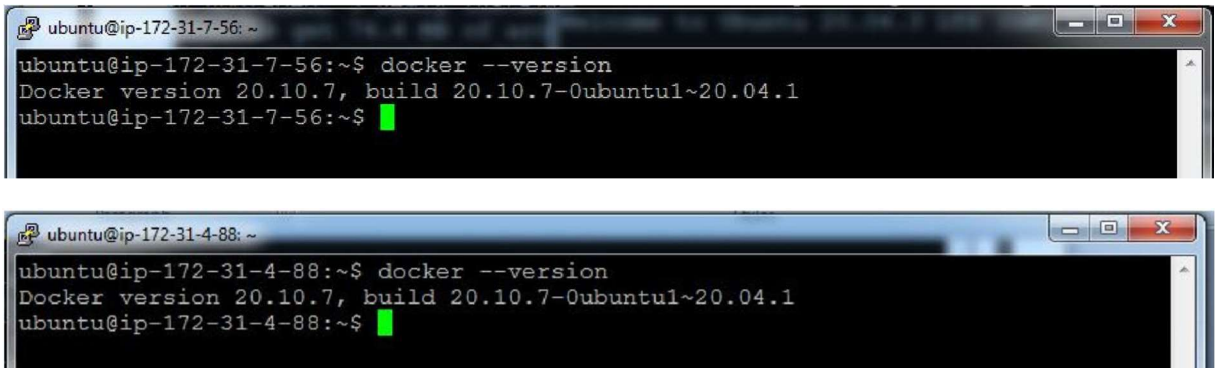

Step 2 : Start and Enable Docker

Set Docker to launch at boot by entering the following:

$ sudo systemctl enable docker

$ sudo systemctl status docker

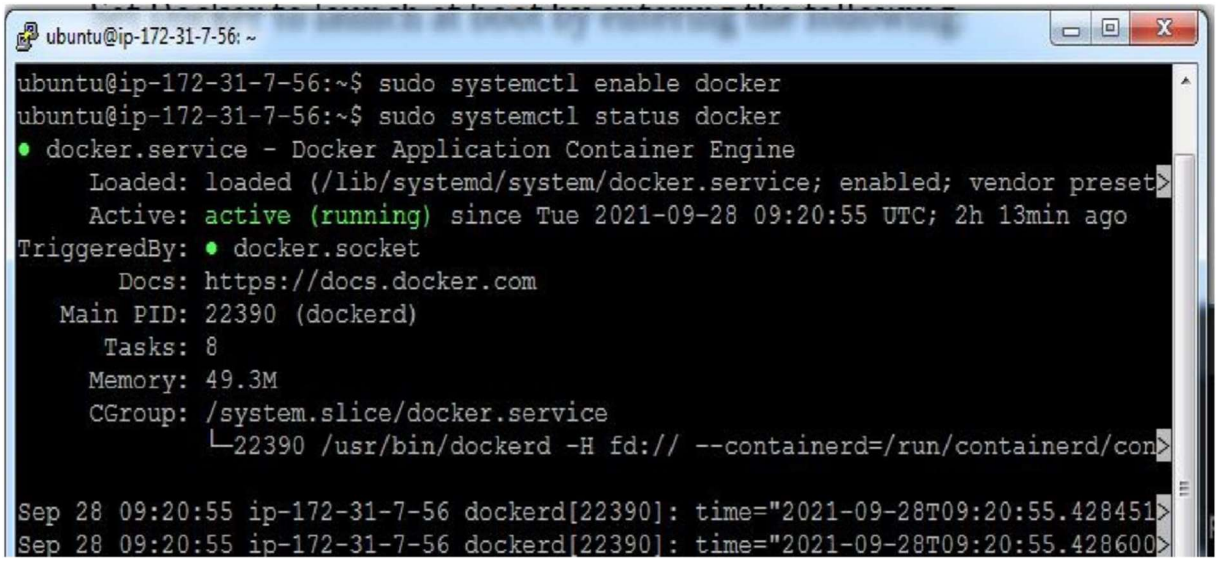

To start Docker if it’s not running:

$ sudo systemctl start docker

Repeat on all the other nodes.

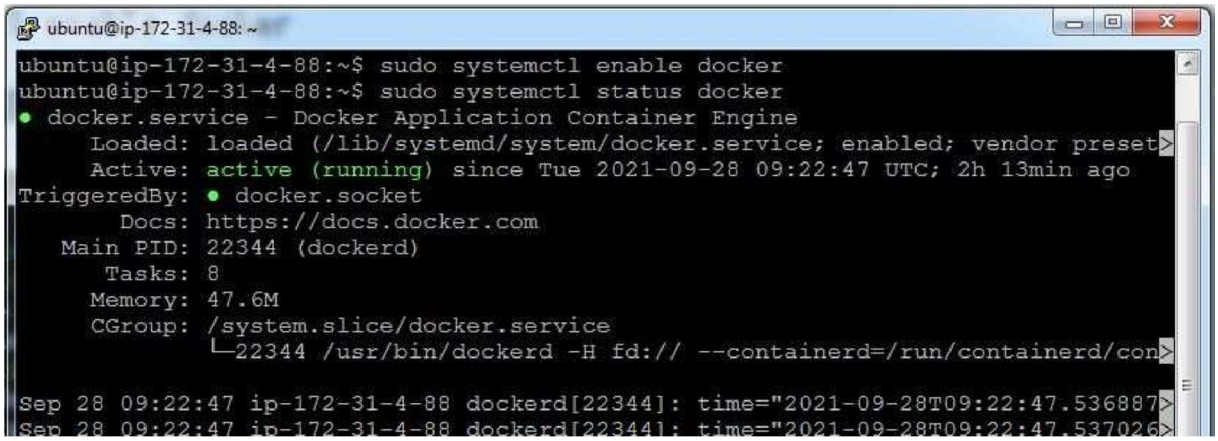

- Install Kubernetes

Step 3 : Add Kubernetes Signing Key

Since you are downloading Kubernetes from a non-standard repository, it is essential to ensure that the software is authentic. This is done by adding a signing key.

Enter the following to add a signing key:

$ curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add

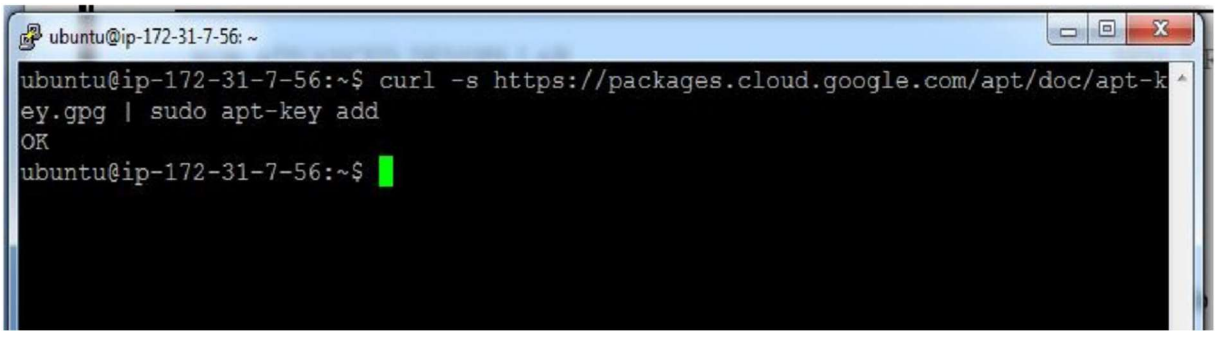

If you get an error that curl is not installed, install it with:

$ sudo apt-get install curlThen repeat the previous command to install the signing keys. Repeat for each server node.

Step 4: Add Software Repositories

Kubernetes is not located in the default repositories. To add them enter the following:

$ sudo apt-add-repository "deb http://apt.kubernetes.io/ kubernetes-xenial main"

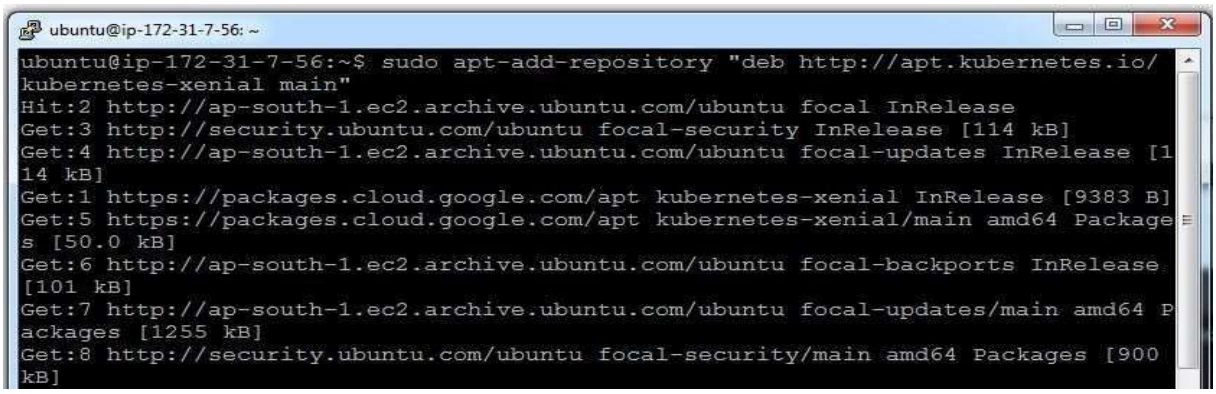

Repeat on each server node.

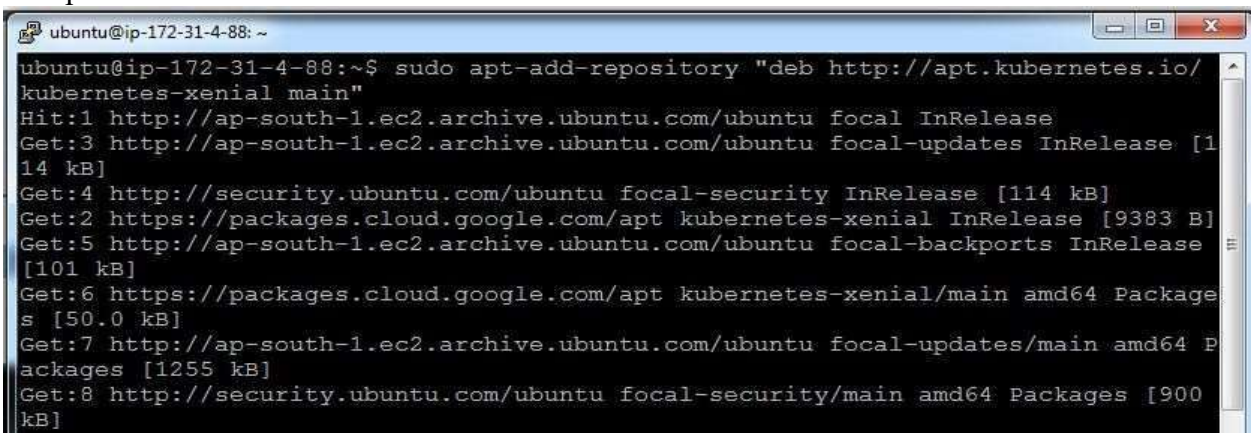

Step 5 : Kubernetes Installation Tools

Kubeadm (Kubernetes Admin) is a tool that helps initialize a cluster. It fast-tracks setup by using community-sourced best practices. Kubelet is the work package, which runs on every node and starts containers. The tool gives you command-line access to clusters.

Install Kubernetes tools with the command:

$ sudo apt-get install kubeadm kubelet kubectl

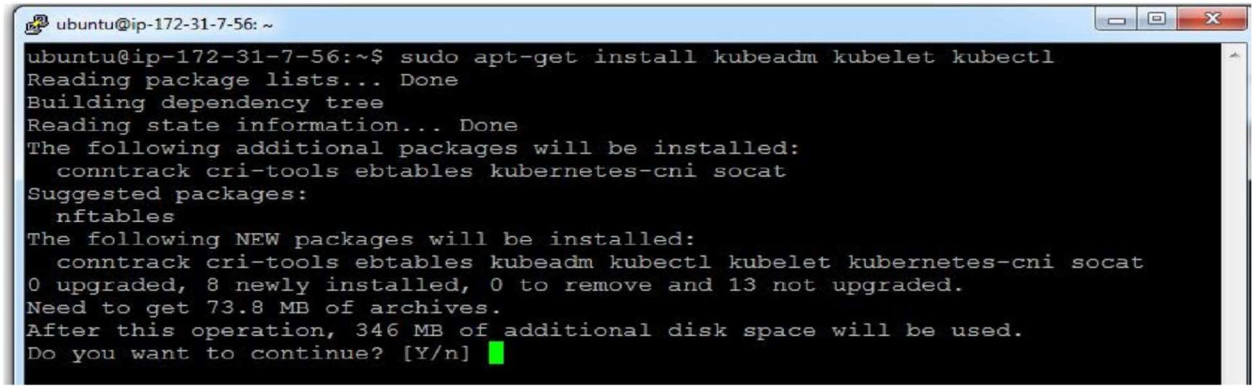

$ sudo apt-mark hold kubeadm kubelet kubectl

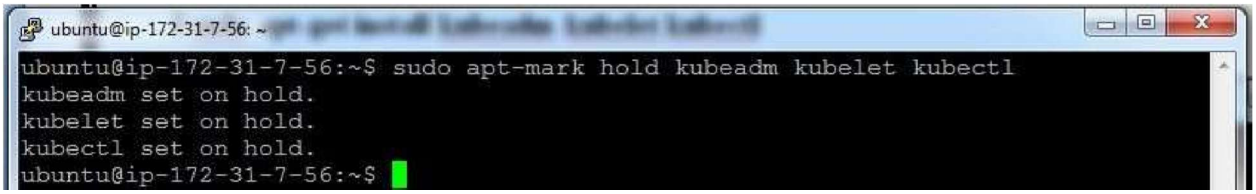

Allow the process to complete. Verify the installation with:

$ kubeadm version

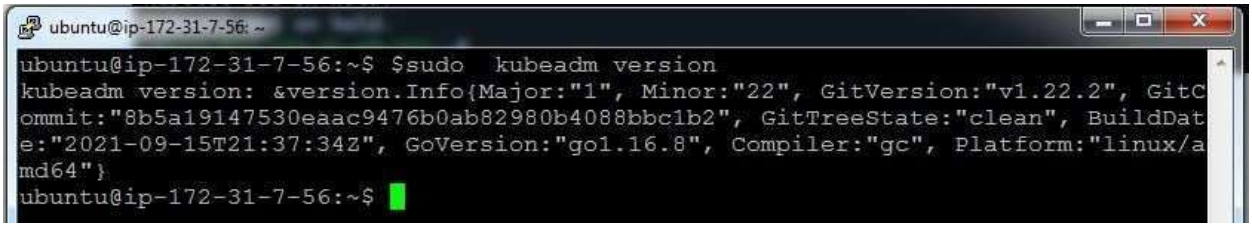

Repeat for each server node.

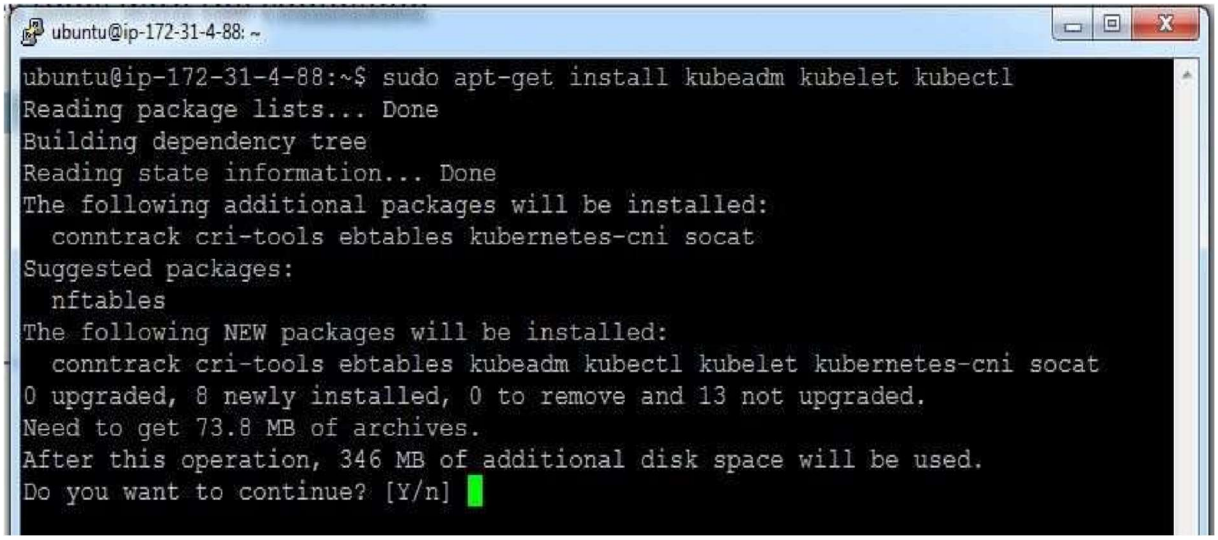

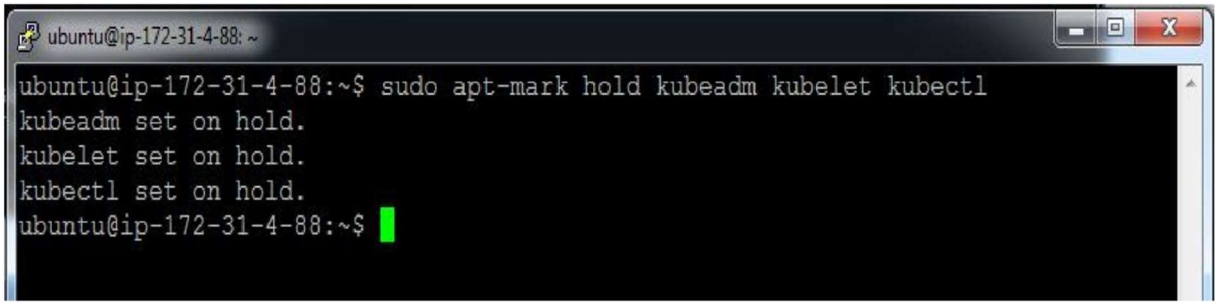

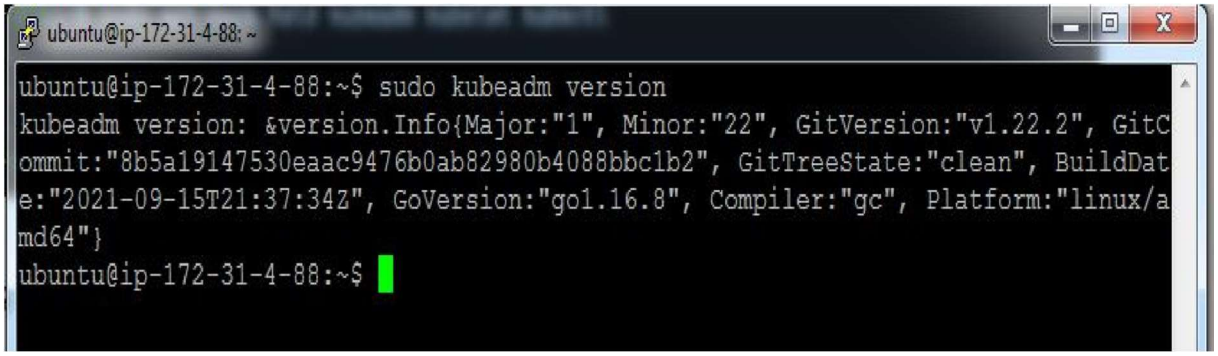

- Kubernetes Deployment

Step 6 : Begin Kubernetes Deployment

Start by disabling the swap memory on each server:

$ sudo swapoff –a

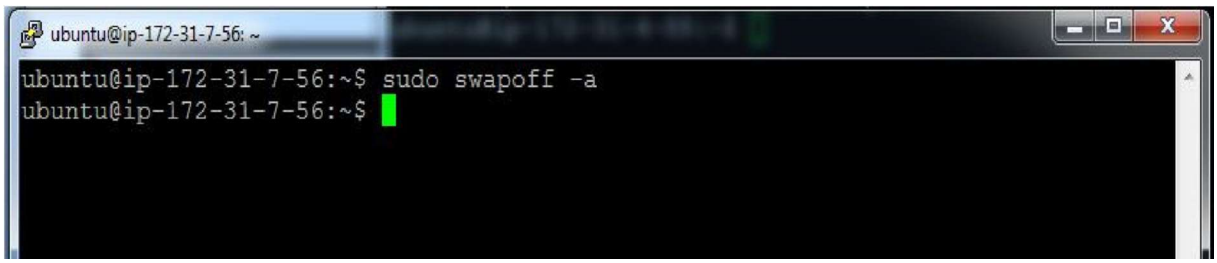

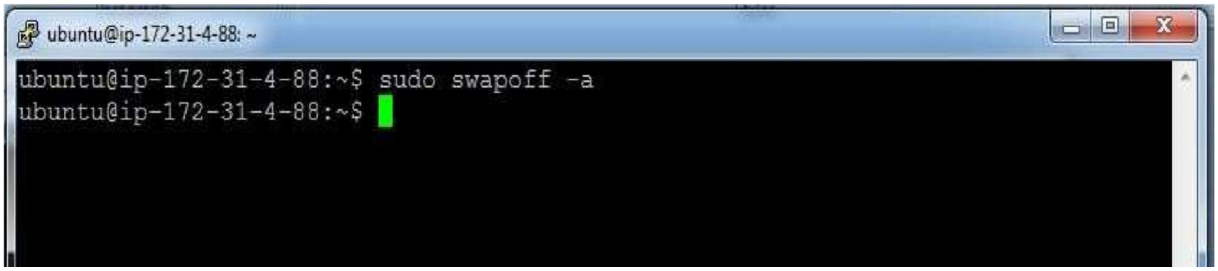

Step 7 : Assign Unique Hostname for Each Server Node

Decide which server to set as the master node. Then enter the command:

$ sudo hostnamectl set-hostname master-node

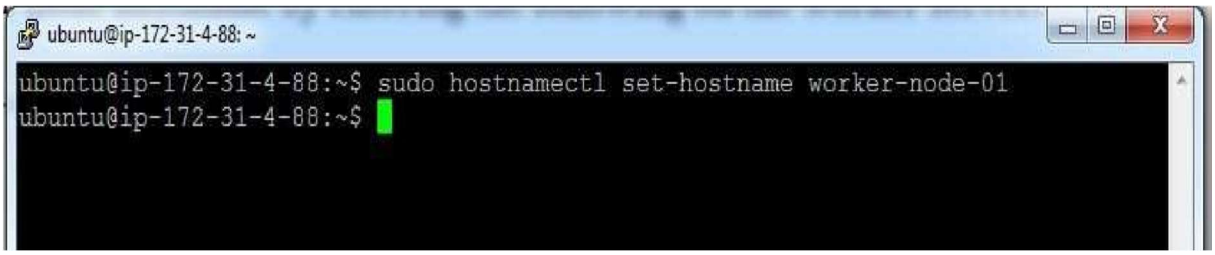

Next, set a worker node hostname by entering the following on the worker server:

$ sudo hostnamectl set-hostname worker-node-01

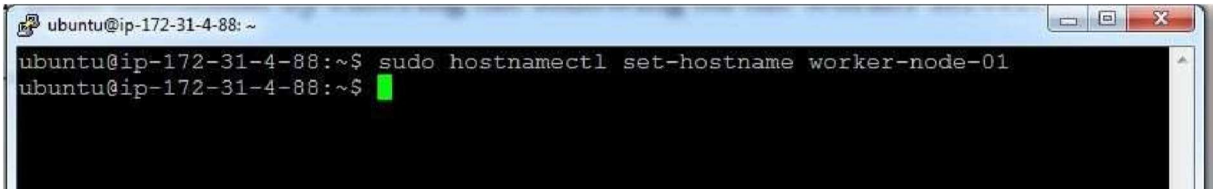

If you have additional worker nodes, use this process to set a unique hostname on each. Logout and again login to see the results.

Step 8 : Initialize Kubernetes on Master Node

Switch to the master server node, and enter the following:

$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16

If you are trying to run this on EC2 you'll get an error message saying less CPU and memory to override the error run the above command with --ignore-preflight-errors=all

For eg:

on-master $ sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=all

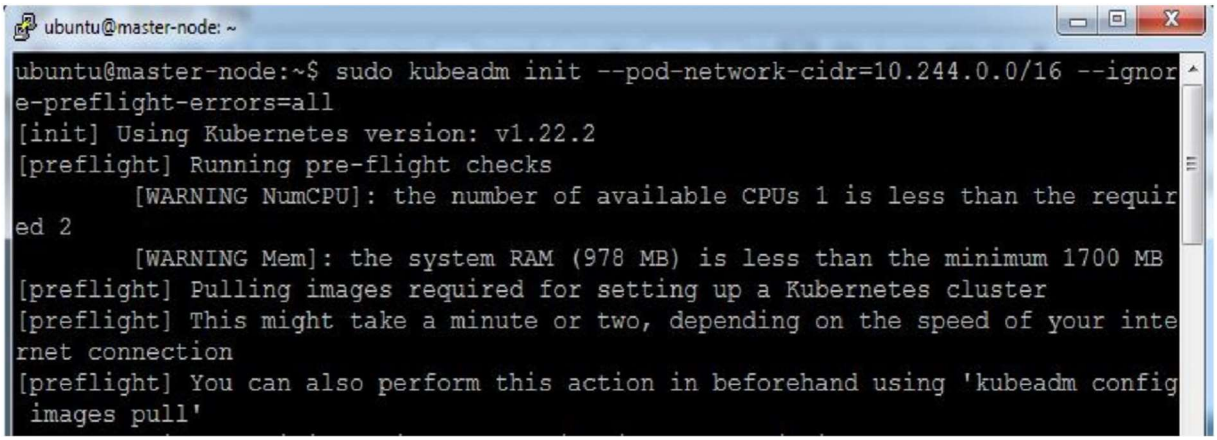

If the kubeadm init command ran without error then ignore this part. If you receive this error "kubelet isn't running or healthy", then do the following.

Create file daemon.json in /etc/docker/ and add following lines in the file.

{

"exec-opts": ["native.cgroupdriver=systemd"]

}

And run the following commands.

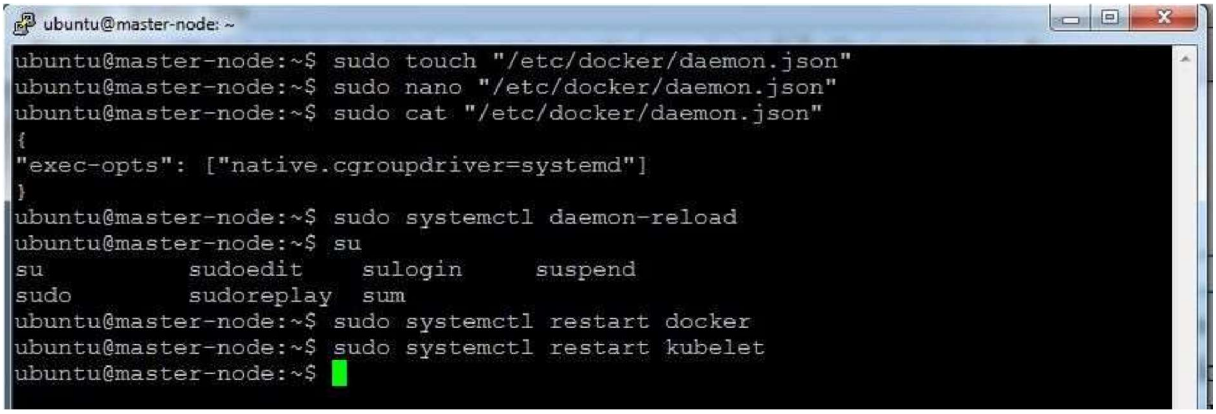

Do this on both master and worker nodes.

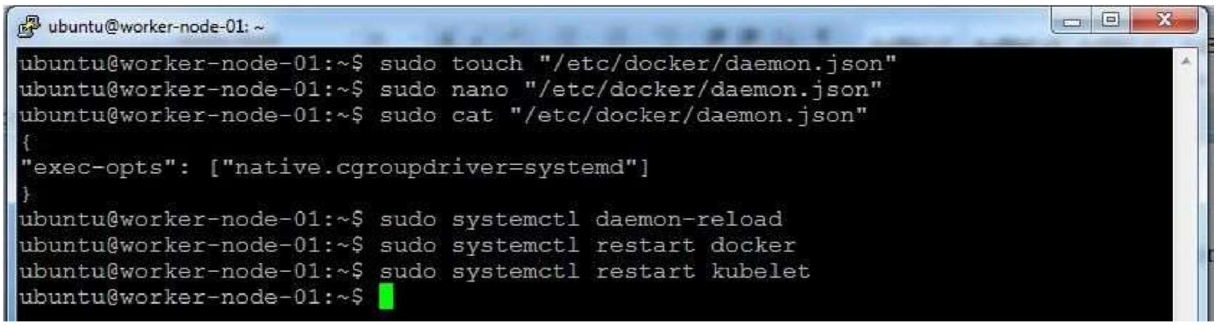

After this run sudo kubeadm reset command and then the init or join command .

Once this command finishes, it will display a kubeadm join message at the end. Make a note of the whole enty. This will be used to join the worker nodes to the cluster.

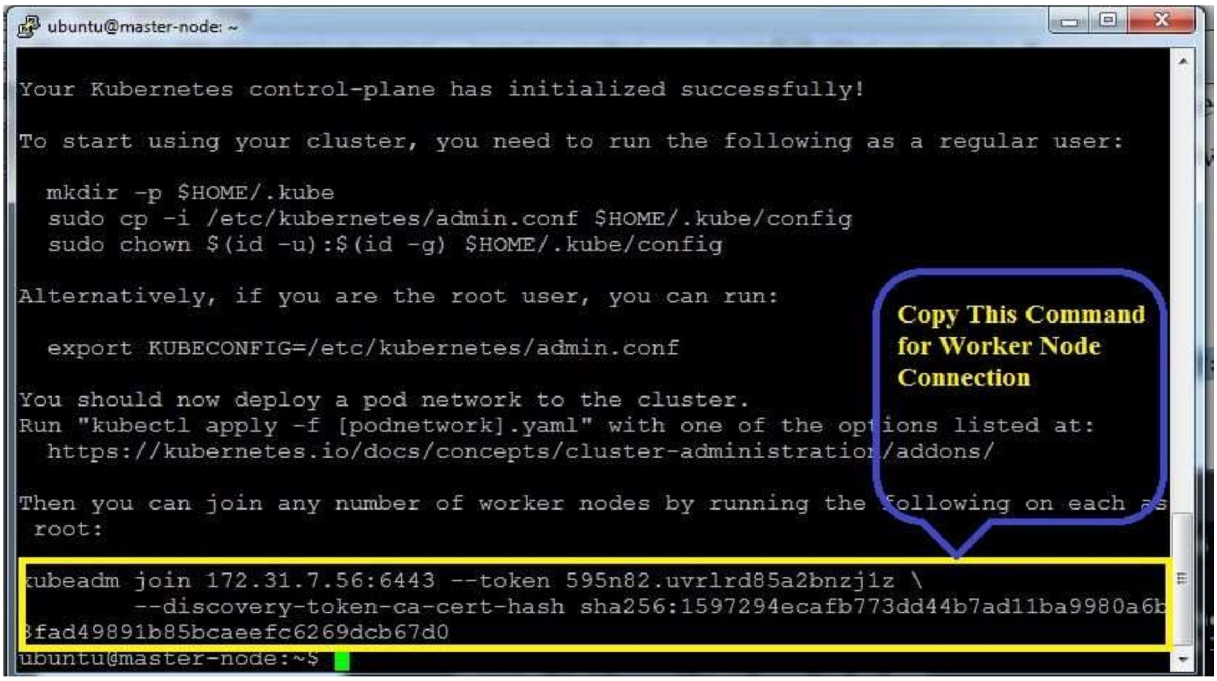

Next, enter the following to create a directory for the cluster:

kubernetes-master $ mkdir -p $HOME/.kube

kubernetes-master $ sudo cp -i /etc/kubernetes/admin.conf

$HOME/.kube/config

kubernetes-master $ sudo chown $(id -u):$(id -g) $HOME/.kube/config

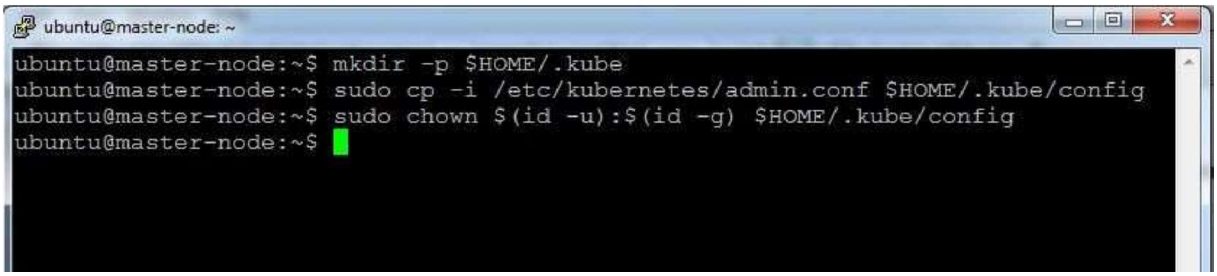

Step 9 : Deploy Pod Network to Cluster

A Pod Network is a way to allow communication between different nodes in the cluster. This tutorial uses the flannel virtual network.

Enter the following:

$ kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kubeflannel.yml

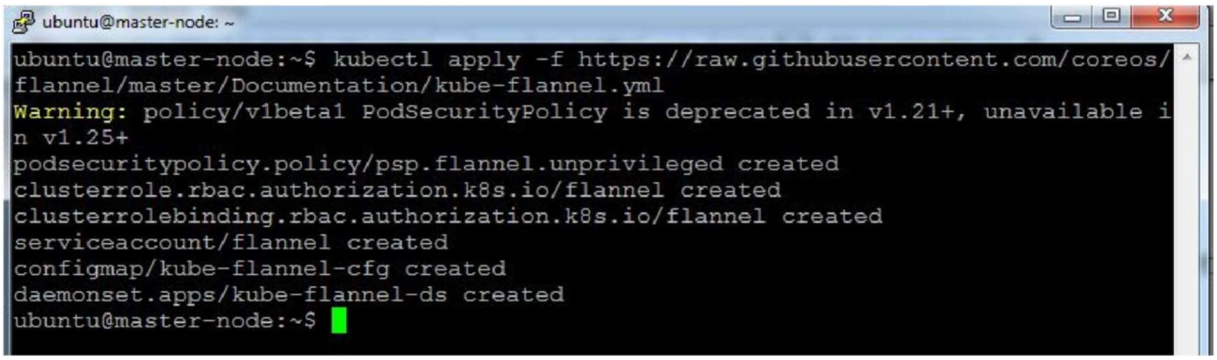

Allow the process to complete.

Verify that everything is running and communicating:

$ kubectl get pods --all-namespaces

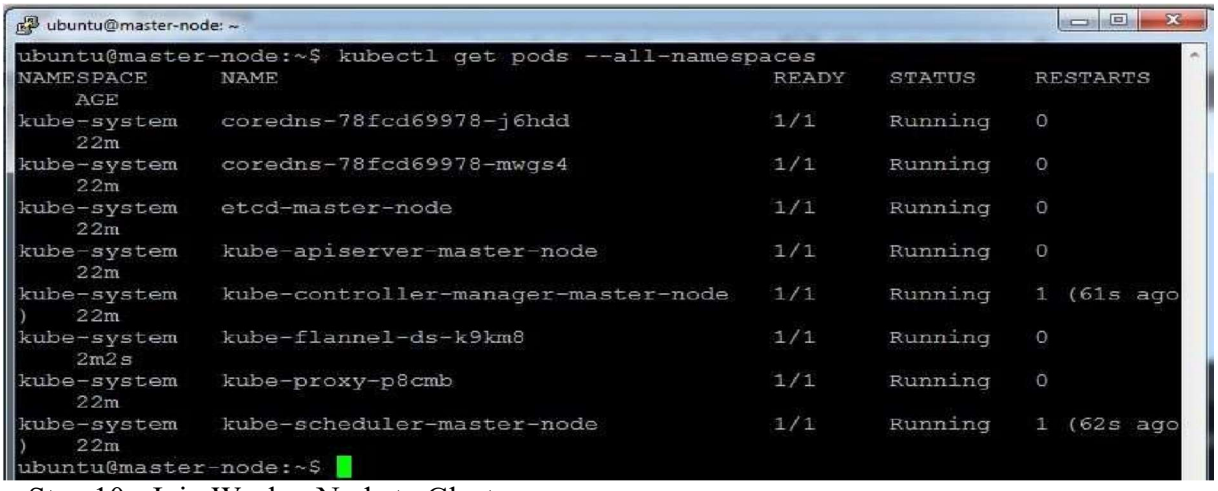

Step 10 : Join Worker Node to Cluster

As indicated in Step 7, you can enter the kubeadm join command on each worker node to connect it to the cluster.

Switch to the worker-node-01 system and enter the command you noted from Step 7:

$ kubeadm join --discovery-token abcdef.1234567890abcdef --discovery-token-ca-cert-hash sha256:1234..cdef 1.2.3.4:6443

Replace the alphanumeric codes with those from your master server. Repeat for each worker node on the cluster. Wait a few minutes; then you can check the status of the nodes.

If you are trying to run this on EC2 you'll get an error message saying less CPU and memory to override the error run the above command with --ignore-preflight-errors=all

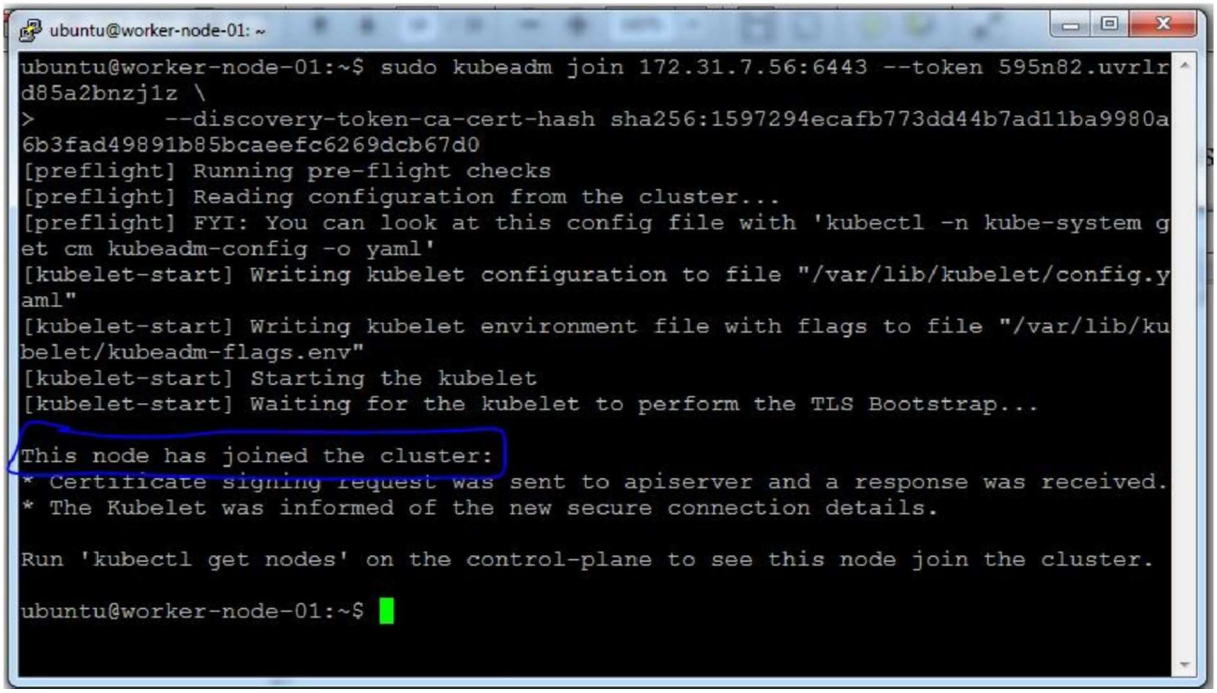

Switch to the master server, and enter:

$ kubectl get nodes

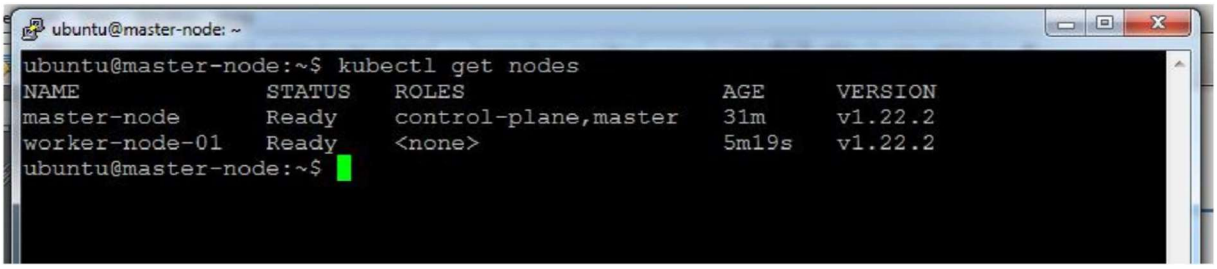

The system should display the worker nodes that you joined to the cluster.
How to gracefully rezove a node from Kubernetes?

On Master Node, find the node

kubernetes-master $ kubectl get nodes

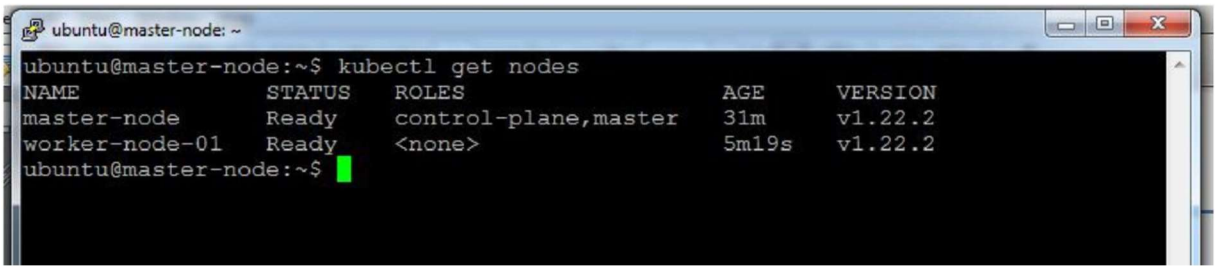

Drain it

kubernetes-master $kubectl drain name_of_node

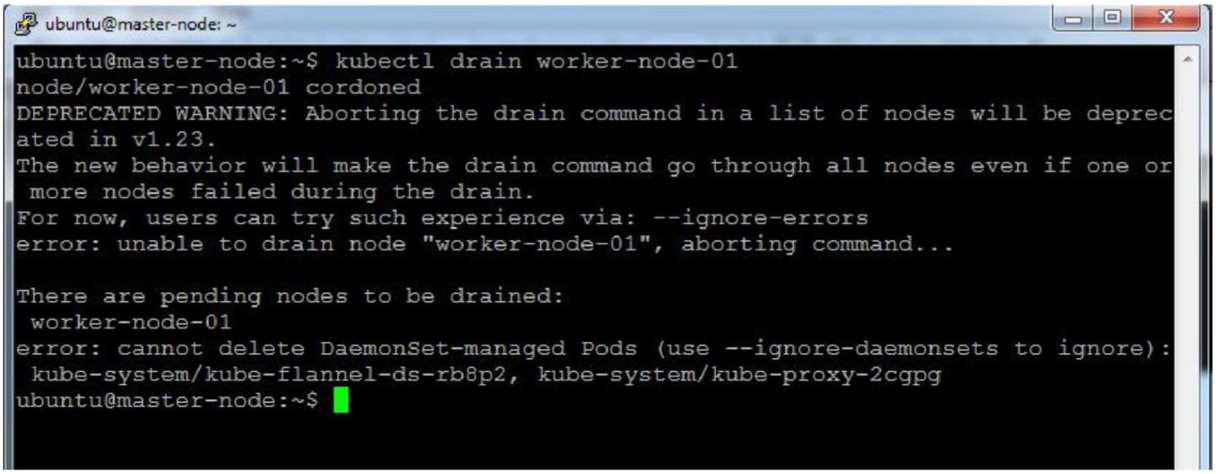

Delete it

kubernetes-master $ kubectl delete node name_of_node

On Worker Node (node-to-be-removed). Remove join/init setting from node

kubernetes-slave $ kubeadm reset

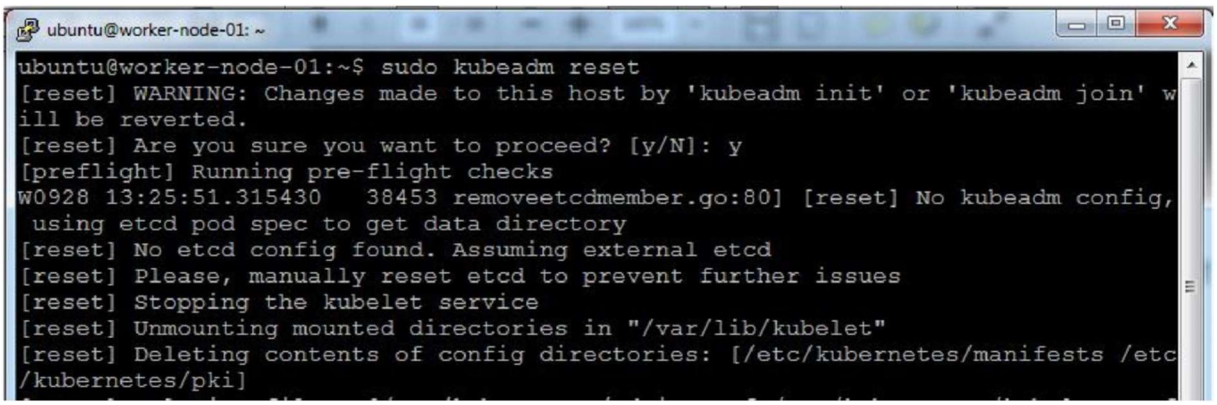

Output: all the containers and service related to the kubernetes cluster are deleted.

kubernetes-master $ kubectl get nodes

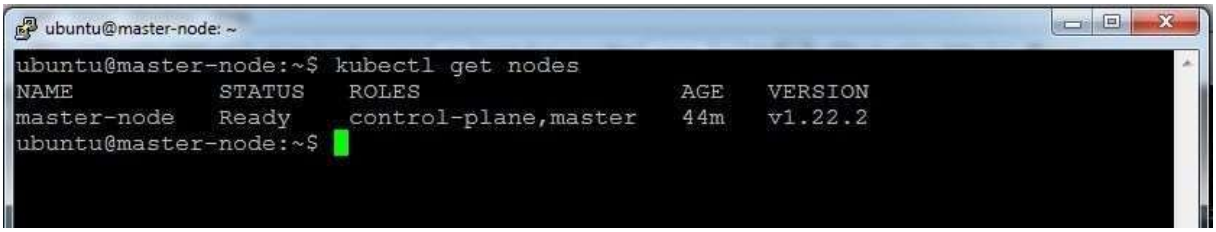

The node has been removed from Kubernetes-master.

Conclusion:

The primary strength of Kubernetes is its modularity and generality. Nearly every kind of application that you might want to deploy you can fit within Kubernetes, and no matter what kind of adjustments or tuning you need to make to your system, they're generally possible. Kubernetes is a production-grade container orchestration system that helps you maximize the benefits of using containers. Kubernetes provides you with a toolbox to automate deploying, scaling, and operating
containerized applications in production. After following all the steps kubernetes successfully.